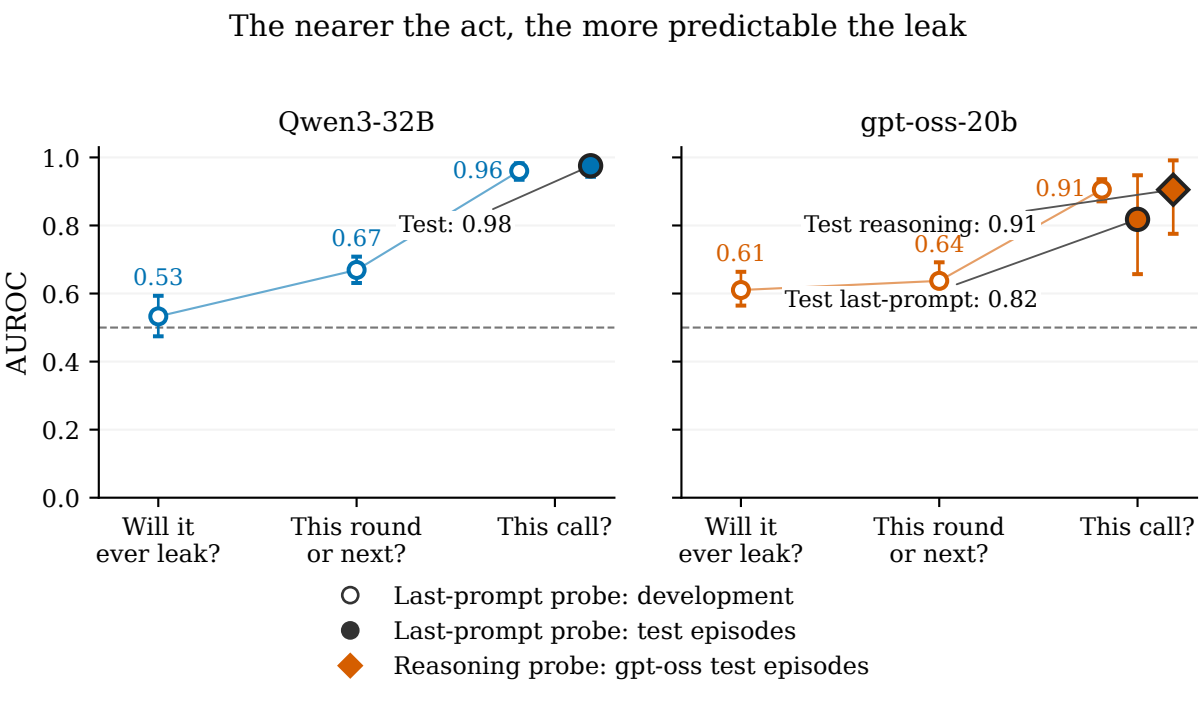
The nearer the act, the more predictable the leak
Qwen3-32B
gpt-oss-20b
AUROC
1.0
0.8
0.6
0.4
0.2
0.0
0.53
0.67
0.96
Test: 0.98
0.61
0.64
0.91
Test reasoning: 0.91
Test last-prompt: 0.82
Will it ever leak?
This round or next?
This call?
Last-prompt probe: development
Last-prompt probe: test episodes
Reasoning probe: gpt-oss test episodes