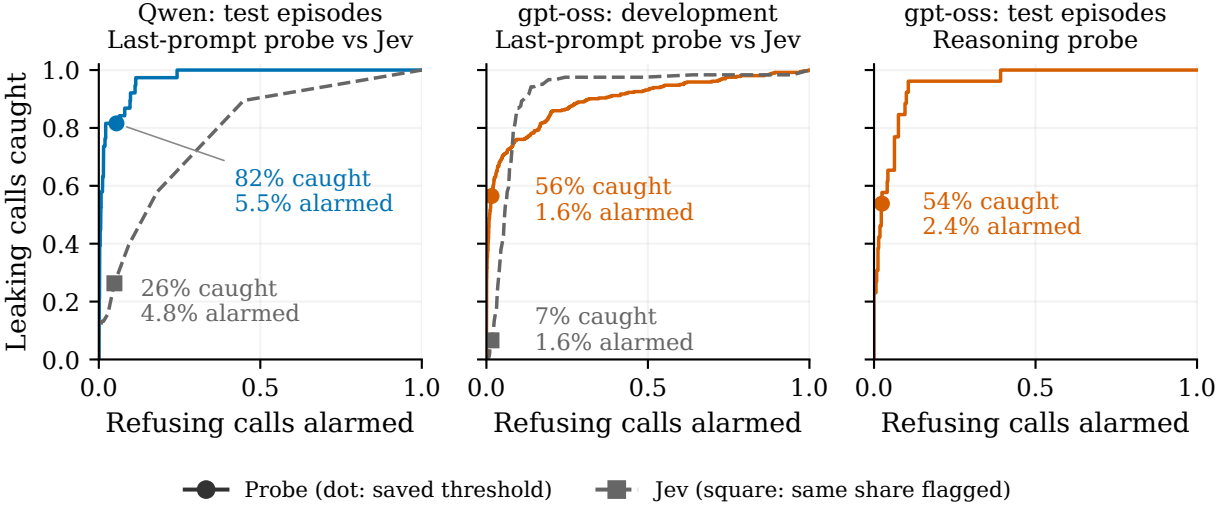
Qwen: test episodes
Last-prompt probe vs Jev
gpt-oss: development
Last-prompt probe vs Jev
gpt-oss: test episodes
Reasoning probe
Leaking calls caught
1.0
0.8
0.6
0.4
0.2
0.0
82% caught
5.5% alarmed
26% caught
4.8% alarmed
56% caught
1.6% alarmed
7% caught
1.6% alarmed
54% caught
2.4% alarmed
0.0
0.5
1.0
Refusing calls alarmed
Probe (dot: saved threshold)
Jev (square: same share flagged)